RcppCNPy: Reading and writing NumPy binary files

Dirk Eddelbuettel

RcppCNPy version 0.2.2 as of November 29, 2013

Abstract

This document introduces the **RcppCNPy** package for reading and writing files created by or for the **NumPy** module for Python.

RcppCNPy is based on **cnpy**, a C++ library written by Carl Rogers.

1 Motivation

Python[1] is a widely-used programming language. It is deployed for use cases ranging from simple scripting to larger-scale application development. Python is also popular for quantitative and scientific application due to the existence of extension modules such as **NumPy**[2] (which is shorthand for Numeric Python).

NumPy can be used for N-dimensional arrays, and provides an efficient binary storage model for these files. In practice, N is often equal to two, and matrices processed or generated in Python can be stored in this form.

R has (as of mid-2012) no dedicated reading or writing functionality for these files. However, Carl Rogers has provided a small C++ library called **cnpy**[3] which is released under the MIT license. Using the 'Rcpp modules' feature in **Rcpp**[4][5], we provide (some) features of this library to R.

2 Examples

2.1 Data creation in Python

The first code example simply creates two files in Python: a two-dimensional rectangular array as well as a vector.

```
>>> import numpy as np
>>>
>>> mat = np.arange(12).reshape(3,4) *
1.1
>>> mat
array([[  0. ,   1.1,   2.2,   3.3],
```

```
       [  4.4,   5.5,   6.6,   7.7],
       [  8.8,   9.9,  11. ,  12.1]])
>>> np.save("fmat.npy", mat)
>>>
>>> vec = np.arange(5) * 1.1
>>> vec
array([ 0. ,  1.1,  2.2,  3.3,  4.4])
>>> np.save("fvec.npy", vec)
>>>
```

As illustrated, Python uses the Fortran convention for storing matrices and higher-dimensional arrays: a matrix constructed from a single sequence has its first consecutive elements in its first row—whereas R, following the C convention, has these first few values in its first column. This shows that to go back and forth we need to transpose these matrices (which represented internally as two-dimensional arrays).

2.2 Data reading in R

We can read the same data in R using the `npyLoad()` function provided by the **RcppCNPy** package:

```
R> library(RcppCNPy)
Loading required package: Rcpp
R>
R> mat <- npyLoad("fmat.npy")
R> mat
     [,1] [,2] [,3] [,4]
[1,]  0.0  1.1  2.2  3.3
[2,]  4.4  5.5  6.6  7.7
[3,]  8.8  9.9 11.0 12.1
R>
R> vec <- npyLoad("fvec.npy")
R> vec
[1] 0.0 1.1 2.2 3.3 4.4
R>
```

The Fortran-order of the matrix is preserved; we obtain the exact same data as we stored.

2.3 Reading compressed data in R

A useful extension to the **cnpy** is the support of **gzip**-compressed data.

---

[1] http://www.python.org
[2] http://numpy.scipy.org/
[3] https://github.com/rogersce/cnpy
[4] Eddelbuettel and François, 2011, JSS, 40(8), http://www.jstatsoft.org/v40/i08/
[5] http://CRAN.R-Project.org/package=Rcpp

```
R> mat2 <- npyLoad("fmat.npy.gz")
R> mat2
     [,1] [,2] [,3] [,4]
[1,]  0.0  1.1  2.2  3.3
[2,]  4.4  5.5  6.6  7.7
[3,]  8.8  9.9 11.0 12.1
```

Support for writing compressed files has been added in version 0.2.0.

2.4 Data writing in R

Matrices and vectors can be written to files using the `npySave()` function.

```
R> set.seed(42)
R> m <- matrix(sort(rnorm(6)), 3, 2)
R> m
           [,1]      [,2]
[1,] -0.564698 0.404268
[2,] -0.106125 0.632863
[3,]  0.363128 1.370958
R> npySave("randmat.npy", m)
R>
R> v <- seq(10, 12)
R> v
[1] 10 11 12
R> npySave("simplevec.npy", v)
```

2.5 Data reading in Python

Reading the data back in Python is straightforward too:

```
>>> m = np.load("randmat.npy")
>>> m
array([[-0.56469817,  0.40426832],
       [-0.10612452,  0.6328626 ],
       [ 0.36312841,  1.37095845]])
>>>
>>> v = np.load("simpleve.npy")
>>> v
array([ 10.,  11.,  12.])
```

3 Performance

The R script `timing` in the `demo/` directory of package **RcppCNPy** provides a simple benchmark. Given two values n and k, a matrix of size $n \times k$ is created with n rows and k columns. It is written to temporary files in i) ascii format using `write.table()`; ii) NumPy format using `npySave()`; and iii) NumPy format using `npySave()` with compression via the `zlib` library (used also by `gzip`).

Table 1 shows some timing comparisons for a matrix with five million elements. Reading the `npy` is clearly

Access method	Time in sec.	Relative to best
npyLoad(pyfile)	1.95	1.00
npyLoad(pygzfile)	4.92	2.53
read.table(txtfile)	128.85	66.24

Table 1: Performance comparison of data reads using a matrix of size $10^5 \times 50$. File size are 39.7mb for ascii, 40.0mb for npy and 10.8mb for npy.gz. Ten replications were performed, and total times are shown.

fastest as it required only parsing of the header, followed by a single large binary read (and the transpose required to translate the representation used by R). The compressed file requires only one-fourth of the disk space, but takes approximately 2.5 times as long to read as the binary stream has be transformed. Lastly, the default ascii reading mode is clearly by far the slowest.

4 Limitations

4.1 Integer support

Support for integer data types is conditional on use of either the `-std=c++0x` or the newer `-std=c++11` compiler extension switches. Only the newer standard supports the `long long int` type needed to represent `int64` data on a 32-bit OS. So until R allows `-std=c++0x` on CRAN packages, users will need to rebuild **RcppCNPy** with the switch enabled. As shown in the previous examples, integers also transparently convert to float types.

4.2 Higher-dimensional arrays

Rcpp supports three-dimensional arrays, this could be support in **RcppCNPy** as well.

4.3 npz files

The **cnpy** library supports reading and writing of sets of arrays; this feature could also be exported.

5 Summary

The **RcppCNPy** package provides simple reading and writing of **NumPy** files, using the **cnpy** library.

Reading of compressed files is also supported as an extension. This offers users a balance between more compact storage at the prices of slightly longer read times.